

00_phase1_log

Phase 1 — Baseline Training Report

Status: COMPLETED

Date: 2026-02-27

Duration: ~83 minutes (10 epochs)

Best Val IoU: **0.2971** (Epoch 10)

1. Objective

Run the provided `train_segmentation.py` script **without any modifications** to establish a baseline performance metric. This gives us a reference point to measure all future improvements against.

2. Environment & Hardware

Component	Details
GPU	NVIDIA GeForce RTX 3050 6GB Laptop GPU
CUDA	12.6
PyTorch	2.10.0+cu126
OS	Windows 10
Python	3.11.9
RAM	16 GB DDR4

3. Training Configuration

Parameter	Value
Backbone	DINOv2 ViT-Small (<code>dinov2_vits14</code>) — frozen
Seg Head	ConvNeXt-style (simple Conv2d stack)
Optimizer	SGD (lr=1e-4, momentum=0.9)

Parameter	Value
Learning Rate	1e-4 (constant, no scheduler)
Epochs	10
Batch Size	2
Image Size	476×266 (nearest multiple of 14)
Loss	CrossEntropyLoss (unweighted)
Augmentations	None
Mixed Precision	No
Checkpointing	Final model only

Why these settings?

The baseline uses the [exact script provided](#) by the hackathon organizers. A DINOv2 ViT-Small backbone is frozen (only the segmentation head is trained). This is intentionally minimal — only 10 epochs, no augmentations, no scheduler.

4. Dataset Split

Split	Samples
Training	2,857 images
Validation	317 images

All images are 960×540 RGB, resized to 476×266 for the ViT-14 patch grid.

5. Per-Epoch Results

Epoch	Train Loss	Val Loss	Train IoU	Val IoU	Train Dice	Val Dice	Train Acc	Val Acc
1	1.1907	0.9890	0.2303	0.2211	0.3256	0.3430	65.28%	65.52%
2	0.9447	0.9104	0.2695	0.2513	0.3775	0.3815	67.51%	67.73%
3	0.8965	0.8788	0.2836	0.2636	0.3962	0.3983	68.37%	68.58%
4	0.8713	0.8605	0.2942	0.2733	0.4076	0.4120	68.90%	69.10%
5	0.8553	0.8457	0.2972	0.2758	0.4116	0.4149	69.29%	69.49%
6	0.8438	0.8353	0.3067	0.2851	0.4241	0.4277	69.59%	69.78%

Epoch	Train Loss	Val Loss	Train IoU	Val IoU	Train Dice	Val Dice	Train Acc	Val Acc
7	0.8352	0.8282	0.3095	0.2874	0.4271	0.4297	69.78%	69.97%
8	0.8286	0.8224	0.3170	0.2924	0.4363	0.4362	69.95%	70.15%
9	0.8228	0.8176	0.3155	0.2934	0.4345	0.4378	70.06%	70.26%
10	0.8184	0.8136	0.3216	0.2971	0.4427	0.4416	70.22%	70.41%

6. Epoch-by-Epoch Analysis

Epoch 1 — Cold Start

- **Val IoU: 0.2211** — The model starts with random segmentation head weights while the DINOv2 backbone provides pre-trained features.
- Loss is high (1.19 train, 0.99 val) as the head hasn't learned to map features to classes yet.
- The large gap between train and val loss is expected during initial epochs.

Epochs 2-4 — Rapid Improvement

- **Val IoU jumps from 0.22 → 0.27** — the segmentation head is quickly learning to associate DINOv2 features with class labels.
- Loss drops significantly (~20% reduction in val loss).
- The model starts learning the easy classes first (Sky is very distinct visually).

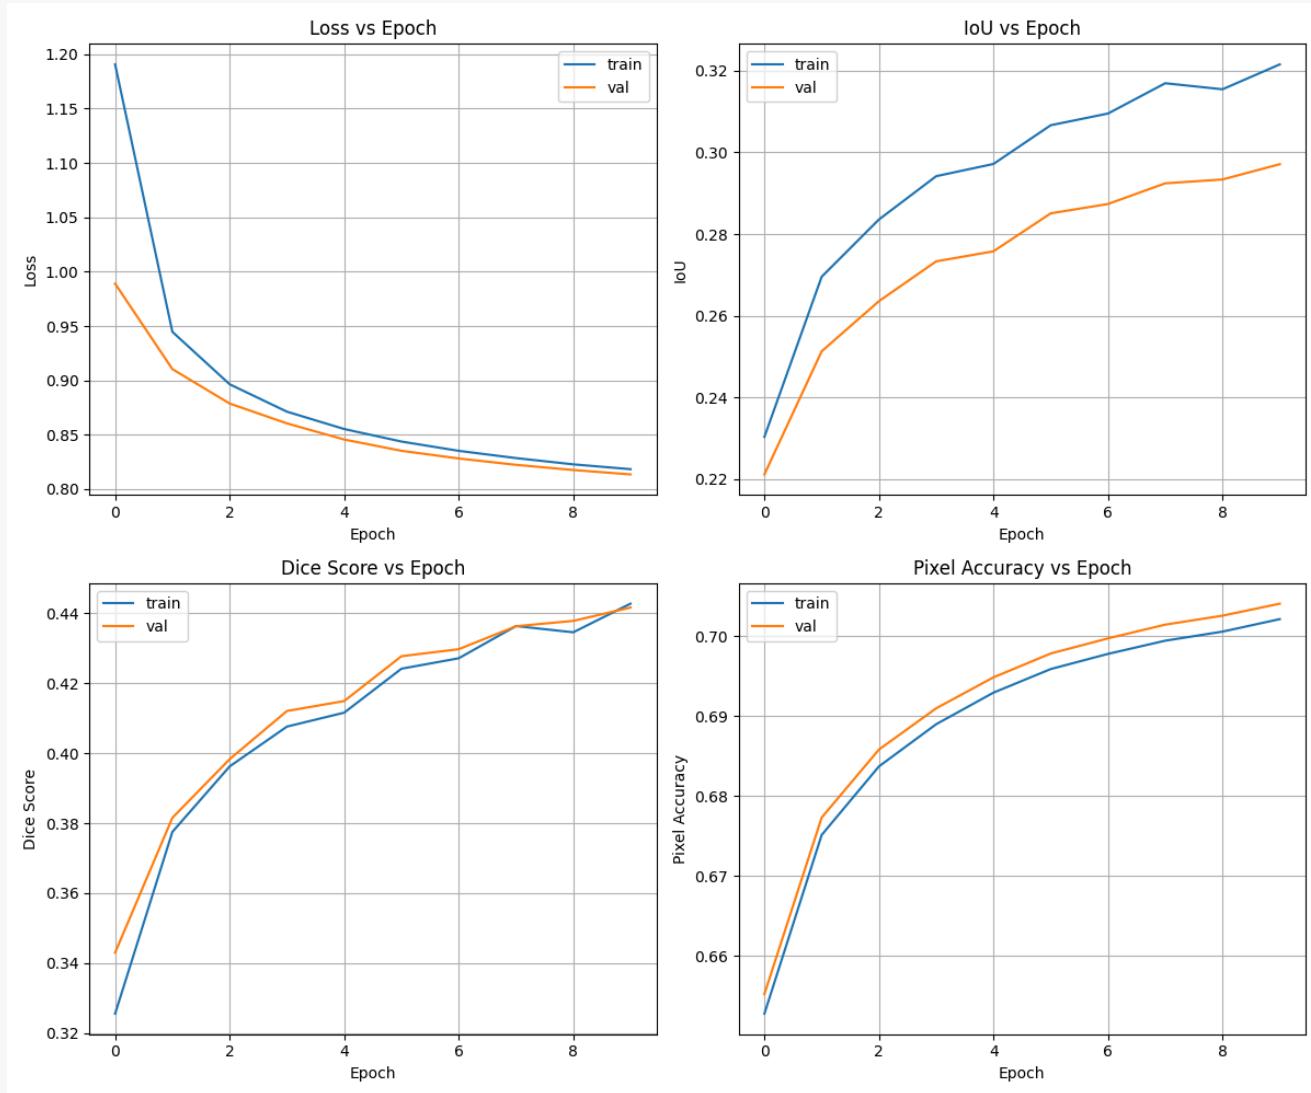
Epochs 5-7 — Slowing Down

- **Val IoU: 0.28 → 0.29** — improvement rate decreases significantly.
- The model has learned the "easy" patterns and now struggles with harder classes (small objects, texturally similar classes).
- Train and val metrics track closely — **no overfitting** detected.

Epochs 8-10 — Still Improving but Slowly

- **Final Val IoU: 0.2971** — model is still improving at epoch 10, suggesting **significant underfitting**.
- Val loss decreased nearly linearly from epoch 1 to 10 — the model hasn't found its minimum yet.
- The gap between train IoU (0.3216) and val IoU (0.2971) is small, confirming **no overfitting**.

7. Training Curve Analysis



What the curves tell us:

- Loss curves (top-left):** Both train and val loss decrease steadily. The near-parallel lines indicate healthy training with no overfitting. Critically, the curves show **no sign of convergence** — they're still going down at epoch 10, meaning more epochs would definitely help.
- IoU curves (top-right):** Steadily increasing but at a decreasing rate. The model went from 0.22 → 0.30 in 10 epochs, but the improvement per epoch got smaller. This is partly because SGD with constant lr=1e-4 is a very slow optimizer for this task.
- Dice curves (bottom-left):** Follow the same trend as IoU. Dice started at 0.33 and ended at 0.44. The Dice-IoU gap is consistent, as expected.
- Accuracy curves (bottom-right):** 65% → 70% over 10 epochs. Note that pixel accuracy can be misleading in segmentation because dominant classes (Sky, Landscape, Dry Grass) make up ~75% of all pixels, so getting them right inflates accuracy while rare classes contribute almost nothing.

8. Key Observations & Diagnosis

✓ What went right:

- **No overfitting** — train and val metrics are nearly identical
- **Stable training** — no spikes, crashes, or NaN values
- **DINOv2 features are useful** — even with a simple head, IoU reaches 0.30

✗ What went wrong:

1. **Severe underfitting** — 10 epochs is far too few. The model was still improving linearly.
2. **No augmentations** — the model sees the exact same images every epoch, limiting generalization.
3. **SGD with constant lr** — very slow convergence. AdamW would be much faster.
4. **No class weighting** — rare classes like Logs (0.07% of pixels) get almost no gradient signal.
5. **Unweighted CrossEntropy** — treats all pixel errors equally, but some classes need more attention.

Root cause of low IoU:

The primary bottleneck is **training time** (only 10 epochs) combined with a **slow optimizer** (SGD at 1e-4). The model has the capacity to do better — it just hasn't trained long enough.

9. Final Scores

Metric	Train	Val
IoU	0.3216	0.2971
Dice	0.4427	0.4416
Pixel Accuracy	70.22%	70.41%
Loss	0.8184	0.8136

10. Output Files

File	Description
evaluation_metrics.txt	Full per-epoch metrics table
all_metrics_curves.png	Combined 2x2 training curves
iou_curves.png	IoU-specific plot
dice_curves.png	Dice-specific plot
training_curves.png	Loss curves

11. Conclusion & Next Steps

Phase 1 establishes a **baseline IoU of 0.2971**. The model is severely undertrained and lacks augmentations, a proper optimizer, and class balancing. The clear path forward is:

1. **More epochs** (30+) to let the model converge
2. **AdamW optimizer** with learning rate scheduler for faster convergence
3. **Data augmentations** to prevent memorization and improve generalization
4. **Class-weighted loss** to handle the extreme class imbalance (Logs = 0.07% vs Sky = 34%)
5. **Mixed precision** to speed up training on the 6GB RTX 3050

These improvements are implemented in **Phase 2**.