

03_Test_Images_and_Scripts_Insights

03 — Test Images & Scripts Insights

Folder Location

```
1  DATASET/
2  └── Offroad_Segmentation_testImages/
3      └── Offroad_Segmentation_testImages/
4          ├── Color_Images/    (342 PNG files)
5          └── Segmentation/   (342 PNG files)
6
7  └── Offroad_Segmentation_Scripts/
8      ├── train_segmentation.py
9      ├── test_segmentation.py
10     ├── visualize.py
11     └── ENV_SETUP/
12         ├── setup_env.bat
13         ├── create_env.bat
14         └── install_packages.bat
```

Test Images

Property	Details
Total Images	342 Color + 342 Segmentation masks
Image Format	.png
Naming Convention	0000060.png to 0000XXX.png — different prefix from train/val (no cc prefix)
Nested Folder	Note the double nesting: Offroad_Segmentation_testImages/Offroad_Segmentation_testImages
Has Segmentation Masks	YES — contrary to the hackathon doc which says "testImages (RGB only, no masks)"

Surprise: The hackathon PDF says test images should be RGB-only with no masks. But your download includes segmentation masks for test images too! This is either a bonus for local evaluation, or the organizers included them by accident. Either way — **NEVER train on test data** (instant disqualification).

Test Dataset vs Train/Val Comparison

Aspect	Train	Val	Test
Count	293	289	342
Filename Prefix	cc	cc	None (just numbers)
Has Masks	Yes	Yes	Yes (unexpected)
Domain	Desert Scene A	Desert Scene A (diff location)	Desert Scene B (different desert)
Purpose	Learn patterns	Monitor overfitting	Final benchmark on unseen locale

Total dataset: 924 images across all three splits.

Scripts Breakdown

train_segmentation.py (591 lines) — The Main Training Script

Config	Value	Notes
Backbone	DINOv2 ViT-Small/14 (dinov2_vits14)	Pre-trained, frozen (no gradients)
Seg Head	SegmentationHeadConvNeXt	Conv stem → depthwise conv block → classifier
Loss	CrossEntropyLoss	Standard for multi-class segmentation
Optimizer	SGD(lr=1e-4, momentum=0.9)	Conservative LR
Batch Size	2	Small — needed for 6GB VRAM
Epochs	10	Very few — likely needs 30-50+
Image Size	476 × 266 (W × H)	Derived from 960/2 and 540/2, snapped to multiples of 14
Classes	10 (including background)	But only 9 annotated + 1 background
Augmentations	NONE	Critical gap — must add augmentations
Scheduler	NONE	No learning rate scheduling

Config	Value	Notes
Best Model Save	NO	Only saves final epoch
Output	train_stats/ folder	Loss curves, IoU curves, Dice curves, metrics text file

Architecture Flow:

- 1 Input Image → DINoV2 (frozen) → Patch Tokens → ConvNeXt Head → Per-pixel logits → Bilinear Upsample → CrossEntropy Loss

Available backbone sizes (in the code but only "small" is used):

- `small` = `vits14` (21M params) ← current
- `base` = `vitb14_reg` (86M params)
- `large` = `vitl14_reg` (300M params)
- `giant` = `vitg14_reg` (1.1B params) — won't fit in 6GB VRAM

test_segmentation.py (488 lines) — Inference & Evaluation Script

Feature	Details
Purpose	Load trained weights → run on test/val images → save predictions + metrics
CLI Arguments	<code>--model_path</code> , <code>--data_dir</code> , <code>--output_dir</code> , <code>--batch_size</code> , <code>--num_samples</code>
Default Model	<code>segmentation_head.pth</code> (in scripts folder)
Default Data	Points to <code>Offroad_Segmentation_testImages/</code>
Outputs	<code>masks/</code> (raw class IDs), <code>masks_color/</code> (RGB visualization), <code>comparisons/</code> (side-by-side), metrics summary
Color Palette	10 defined colors for visualization (black, forest green, lime, tan, brown, olive, saddle brown, gray, sienna, sky blue)

Color Palette for Classes:

Class Color
--- ---
Background Black [0,0,0]
Trees Forest Green [34,139,34]

```
| Lush Bushes | Lime [0,255,0] |
| Dry Grass | Tan [210,180,140] |
| Dry Bushes | Brown [139,90,43] |
| Ground Clutter | Olive [128,128,0] |
| Logs | Saddle Brown [139,69,19] |
| Rocks | Gray [128,128,128] |
| Landscape | Sienna [160,82,45] |
| Sky | Sky Blue [135,206,235] |
```

visualize.py (54 lines) — Quick Mask Colorizer

- Takes a folder of segmentation masks → assigns random colors to unique pixel values
→ saves colorized PNGs
 - **Input folder is blank (" ")** — you need to set it manually
 - Useful for quickly eyeballing what the masks look like
-

ENV_SETUP Scripts (Windows .bat files)

File	What It Does
setup_env.bat	Runs <code>create_env.bat</code> then <code>install_packages.bat</code> in sequence
create_env.bat	<code>conda create --name EDU python=3.10 -y</code>
install_packages.bat	<code>conda install pytorch torchvision pytorch-cuda=11.8 ultralytics -y && pip install opencv-contrib-python && pip install tqdm</code>

Note: These scripts use **Conda** and create an `EDU` environment. Since we already set up a **venv**, we'll install the packages directly into our venv instead, skipping these bat files entirely.

Critical Improvement Opportunities

1. **Add data augmentations** — #1 priority (random flip, rotation, color jitter, random crop)
2. **Increase epochs** — 10 is way too few; try 30–50

3. **Add learning rate scheduler** — `CosineAnnealingLR` or `StepLR`
4. **Save best model** — track `best_val_iou` and save checkpoint
5. **Try larger backbone** — `vitb14_reg` (base) if VRAM allows
6. **Add early stopping** — prevent wasting time if model plateaus
7. **Consider mixed precision** — `torch.cuda.amp` for faster training on RTX 3050
8. **Add class weights to loss** — handle class imbalance (Landscape/Sky dominate)