

MODEL-README

Trained Model Weights

All trained model checkpoints from our 3-phase training journey.

Files

| File | Phase | Val IoU | Size | Description |
|---------------------------------|---------|---------------|-------|---|
| phase2_best_model_iou0.4036.pth | Phase 2 | 0.4036 | 28MB | Best checkpoint from Phase 2 (epoch 26). Contains model state dict + optimizer state + config. Uses DINoV2 ViT-Small backbone + ConvNeXt head. |
| phase3_best_model_iou0.5161.pth | Phase 3 | 0.5161 | 39MB | ★ Best overall model. Best checkpoint from Phase 3 (epoch 40). Contains model state dict + optimizer state + config. Uses DINoV2 ViT-Base backbone + UPerNet head. |
| segmentation_head.pth | Phase 3 | 0.5161 | 9.3MB | Head-only weights extracted from Phase 3 best checkpoint. This is what <code>test_segmentation.py</code> loads for inference. Smaller file because it only contains UPerNet head parameters (no optimizer state). |

How to Use

For Inference (test images)

```
# test_segmentation.py loads segmentation_head.pth automatically
python TRAINING_SCRIPTS/test_segmentation.py
```

For Resuming Training

```
checkpoint = torch.load('MODELS/phase3_best_model_iou0.5161.pth')
model.load_state_dict(checkpoint['model_state_dict'])
optimizer.load_state_dict(checkpoint['optimizer_state_dict'])
print(f'Resuming from epoch {checkpoint["epoch"]}, IoU={checkpoint["val_iou"]:.4f}')
```

Why No Phase 1 Model?

Phase 1's baseline script only saved the final epoch weights (no best checkpoint). The final model (IoU=0.2971) has been superseded by Phase 2 and 3, so it isn't preserved here.